

THE EXPERT'S VOICE® IN ORACLE

# Pro Oracle SQL Development

Best Practices for Writing  
Advanced Queries

—  
Jon Heller

Apress®

⟨IOUG⟩  
independent oracle users group

**You are viewing a courtesy copy of chapter  
16 from Jon Heller's book, *Pro Oracle SQL  
Development***

**Would you like a copy of the complete  
book?**

**Purchase now from Amazon:**

**[https://www.amazon.com/Pro-Oracle-SQL-  
Development-Practices/dp/1484245164/](https://www.amazon.com/Pro-Oracle-SQL-Development-Practices/dp/1484245164/)**

**Or from Apress.com:**

**[https://www.apress.com/us/book/97814842  
45163](https://www.apress.com/us/book/9781484245163)**

**Thank you for your support of the author  
and his work**

# Understand SQL Performance with Algorithm Analysis

Solving Oracle SQL performance issues is the pinnacle of SQL development. Performance tuning requires a combination of all the skills previously discussed in this book. We need to understand the development process (to know why problems happened and weren't caught sooner), advanced features (to find alternative ways to implement code), and programming styles (in order to understand the code and rewrite it into something better).

Programming is one of the few fields with orders of magnitude difference in skill level between professionals. We've all had to deal with coworkers who are only one tenth as productive as us, and we've all been humbled by developers whose code is ten times better than ours. With performance tuning those numbers get even higher, and the rewards are greater. It's exhilarating when we make a small tweak and something runs a million times faster.

But "a million times faster" is only the beginning of this story. All SQL tuning guides discuss run *times*, but none of them consider the run time *complexity*. Database performance tuning requires more than an encyclopedic knowledge of obscure features and arcane settings. Performance tuning calls for a different mindset.

Algorithm analysis is the most underused approach to understanding Oracle performance problems. This chapter tells the story of Oracle performance through the lens of simple algorithm analysis. The sections are not in a traditional order, such as ordered by performance concepts, tuning tools, or performance problem categories. The sections are ordered by time complexity, from fastest to slowest.

Admittedly, using algorithm analysis to understand Oracle operations is not the most useful performance tuning technique. But algorithm analysis should not be relegated to the halls of academia, it should be a part of everyone's SQL tuning toolkit. We'll solve more problems with techniques like sampling and cardinality estimates, but we'll never truly understand performance without understanding the algorithms.

Algorithm analysis can help us with both proactive and reactive performance tuning. For proactive performance tuning we need to be aware of the advantages available if we create different data structures. For reactive performance tuning we need to be able to measure the algorithms chosen by the optimizer and ensure Oracle made the right choices.

The two chapters after this one provide a more traditional approach to performance tuning; those chapters describe the different concepts and solutions used in Oracle SQL tuning. This chapter helps us comprehend precisely why the stakes are so high for the decisions Oracle must make. This material should be useful to SQL developers of any skill level.

# Algorithm Analysis Introduction

With a few simple mathematical functions we can gain a deeper understanding of the decisions and trade-offs involved in execution plan creation. Performance results are often given as simple numbers or ratios, such "X runs in 5 seconds, and Y runs in 10 seconds." The wall-clock time is important, but it's more powerful to understand and explain our results with mathematical functions.

Algorithm analysis, also known as asymptotic analysis, finds a function that defines the boundary of the performance of something. This technique can apply to memory, storage, and other resources, but for our purposes we only need to consider the number of steps in an algorithm. The number of steps is correlated with run time and is an oversimplification of the overall resource utilization of a database system. This chapter ignores measuring different kinds of resource utilization and considers all database "work" to be equal.

A full explanation of algorithm analysis would require many precisely defined mathematical terms. Don't worry, we don't need a college course in computer science to use this approach. A simplified version of algorithm analysis is easily applied to many Oracle operations.

Let's start with a simple, naïve way to search a database table. If we want to find a single value in a table, the simplest search technique would be to check every

row. Let's say the table has  $N$  rows. If we're lucky, we'll find the value after reading only 1 row. If we're unlucky, we have to read  $N$  rows. On average, we will have to read  $N/2$  rows. As the number of table rows grows, the average number of reads grows linearly.

Our naïve algorithm for searching a table has a best-case, average-case, and worst-case run time. If we plot the input size and the number of reads, we'll see that the worst-case performance is *bounded* by the function  $N$ . In practice, we only care about the upper bound. Worst-case run time can be labeled with the traditional Big O notation as  $O(N)$ . A real-world solution would also include some constant amount of work for initialization and cleanup. Constants are important but we can still meaningfully compare algorithms without constants.

Figure 16-1 below visualizes the worst-case as a solid line, an asymptote that the real values can never exceed. Our not-so-smart search algorithm uses the dashed line, and it takes more or less steps<sup>1</sup> depending on exactly where the value is placed in the table. Our real-world results may look like the dashed line, but to understand performance it helps to think about results as the solid line.

---

<sup>1</sup> This book uses the word “steps” to refer to the work performed by algorithms. Traditionally, that value is named “operations”. But Oracle execution plans already prominently use the word “operation”, so it would be confusing if I used the standard name. This chapter isn't a mathematical proof so it doesn't matter if our terminology is non-standard.

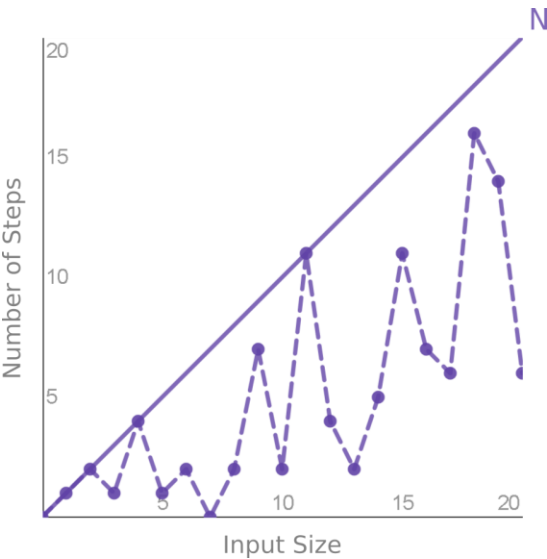


Figure 16-1. Number of steps versus input size for linear search algorithm, including an asymptote.

The functions by themselves are meaningless, it's the comparison between the functions that matter. We want Oracle to choose the functions that have the lowest value on the Y axis, even as the data increases along the X axis.

If Oracle only had to compare straight lines that begin at the origin, then the task would be trivial. The task becomes difficult when each algorithm's performance is represented by a different curve, and the curves intersect at different points along the X axis. Stated as an abstract math problem: to know which curve has the lowest Y value Oracle has to determine the X value. Stated as a practical database problem: to know which operation is fastest Oracle must accurately estimate the number of rows. Bad execution plans occur when Oracle doesn't have the right information to accurately estimate which algorithm is cheaper.

Figure 16-2 below shows the most important functions discussed in this chapter.  $O(1)$ ,  $O(\infty)$ , and Amdahl's law didn't fit together on the graph, but they are also discussed later. Luckily, the most important Oracle operations fall into a small number of categories. We don't need to look at the source code or write proofs, most database operations are easy to classify. Spend a few moments looking at the lines and curves in Figure 16-2 below.

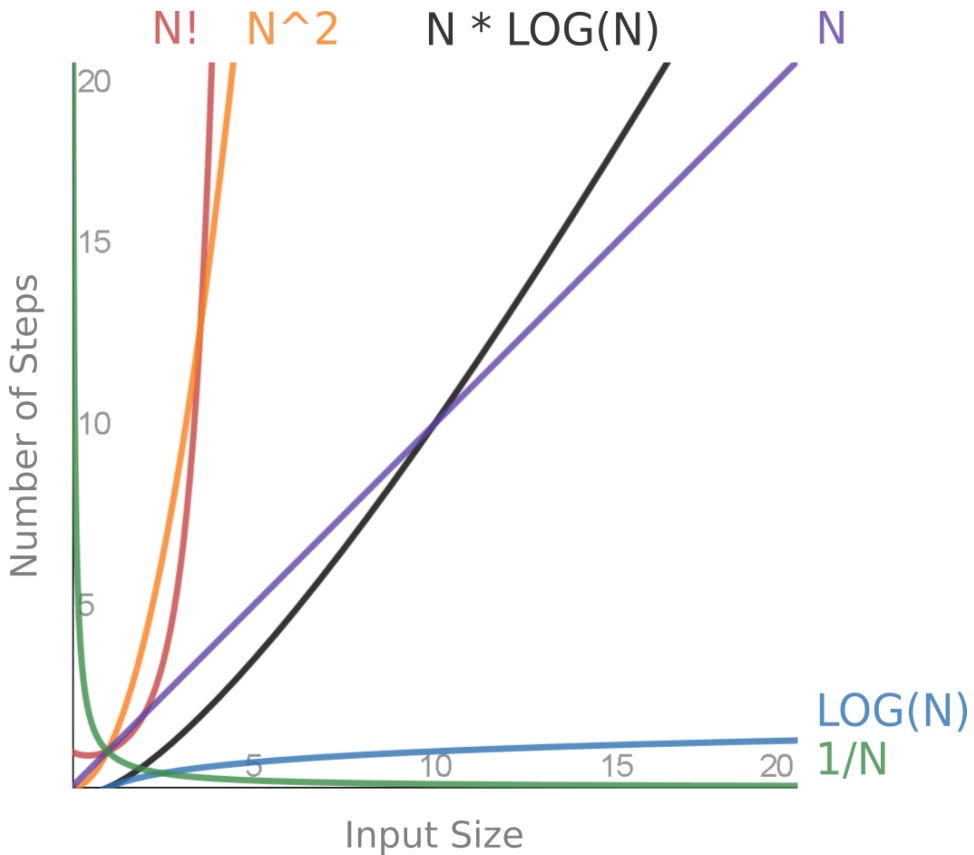


Figure 16-2. Functions that represent important Oracle operations.

The next sections discuss each function, where we find them, and why they matter. Comparing the above shapes helps explain many Oracle performance problems.

## $O(1/N)$ – Batching to Reduce Overhead

The harmonic progression  $1/N$  perfectly describes the overhead of many Oracle operations. This function describes sequence caching, bulk collect limit, prefetch, arraysize, the number of subqueries in a `UNION ALL`, PL/SQL packages with bulk options like `DBMS_OUTPUT` and `DBMS_SCHEDULER`, and many other operations with a configurable batch size. Choosing the right batch size is vital to performance. For example, performance will be disastrous if there is a SQL-to-PL/SQL context switch for every row.

If we want to avoid overhead we have to combine operations, but how many operations do we combine? This is a difficult question with many trade-offs. Not aggregating *anything* is ridiculously slow, but aggregating *everything* will cause memory or parsing problems. Batching is one of the keys to good performance, so we need to think clearly about how much to batch.

Just pick 100 and stop worrying about it. We can be confident in choosing a value like 100 by understanding the charts in Figure 16-3 below. The first thing to notice is that the theoretical results almost perfectly match the real-world results.




Figure 16-3. The effect of increased batch size on total run time.

For the theoretical chart on the left, the total run time is the dotted line of the “real work” plus the solid line of the overhead. By increasing the batch size we can significantly reduce the overhead. But there is only so much overhead to reduce. No matter how much we increase the batch size the line will never get to zero. Our processes are never purely overhead, there should always be a significant amount of real work to do.

The lines in the above graphs start near infinity and quickly plateau. As we increase the batch size the overhead rapidly disappears. A batch size of 2 removes 50% of the overhead, a batch size of 10 removes 90%, and a batch size 100 removes 99%. A setting of 100 is already 99% optimized. The difference between a batch size of 100 and a billion cannot theoretically be more than 1%. And the real-world results almost perfectly match the theory.



The scripts I used to produce the real-world results can be found in the GitHub repository. There’s not enough space to list all the code here but the code is worth briefly discussing. The scripts created two scenarios with a ridiculous amount of overhead; inserting rows where every value is a sequence, and bulk collecting a small amount of data and doing nothing with the results. If any test case is going to show a *meaningful* improvement created by increasing batch size from 100 to 1000, this was it.

Table 16-1 below describes three batching scenarios in detail. These rows are three completely unrelated tasks that can be perfectly explained with the same time complexity.

Table 16-1. Different tasks where performance depends on reducing overhead.

| Task                      | Real Work      | Overhead                      | Configurable Parameter |
|---------------------------|----------------|-------------------------------|------------------------|
| Bulk collect              | Selecting data | SQL and PL/SQL context switch | LIMIT                  |
| INSERT using sequence     | Inserting data | Generating sequence numbers   | CACHE SIZE             |
| Application fetching rows | Selecting data | Network lag                   | Fetch size             |

There are many practical results we can derive from this theory. The default bulk collect size of 100 used by cursor `FOR` loops is good enough; there’s almost never any significant benefit to using a ridiculously-high custom limit. The default sequence caching size of 20 is good enough; in extreme cases it might be worth increasing the cache size slightly, but not by a huge amount. Application prefetch is different for each application; we should aim for something close to the number 100. These rules apply to any overhead-reducing optimization.

If we have an extreme case, or drill way down into the results, we can always find a difference with a larger batch size. But if we find ourselves in a situation where setting a huge batch size helps then we’ve already lost; we need to bring our algorithms to our data, not our data to our algorithms.

In practice, bringing algorithms to our data means we should put our logic in SQL, instead of loading billions of rows into a procedural language for a trivial

amount of processing. If we foolishly load all the rows from a table into PL/SQL just to count them with `V_COUNT := V_COUNT+1`, then increasing the batch size will help. But the better solution would be to use `COUNT(*)` in SQL. If we load a billion rows into PL/SQL and perform real work with those rows, the few seconds we save from a larger batch size will be irrelevant. There are always exceptions, like if we have to access lots of data over a database link with horrendous network lag, but we should not let those exceptions dictate our standards.

There are trade-offs between space and run time, but with a harmonic progression time complexity we will quickly trade space for *no* run time. Developers waste a lot of time debating and tweaking large numbers on the right side of the above graphs, but almost all the bang for our buck happens quickly on the left side of the graph. When we have a  $1/N$  time complexity the point of diminishing returns is reached very quickly. We should spend our time looking for opportunities to batch commands and reduce overhead, not worrying about the precise batch size.

## O(1) – Hashing, Other Operations

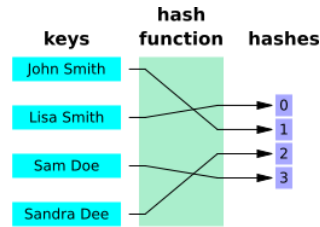
Constant time access is ideal but often unrealistic. The operations that can work in constant time are mostly trivial, such as using a predicate like `ROWNUM = 1`. The constant time function is simple, just a horizontal line, and is not worth showing on a graph. The most important Oracle operation that can run in constant time is hashing. Hashing is a core database operation and is worth discussing in detail.

### How hashing works

Hashing assigns a set of items into a set of hash buckets. Hash functions can assign items into a huge number of cryptographically random buckets, such as `STANDARD_HASH('some_string', 'SHA256')`. Or hash functions can assign values into a small number of predefined buckets, such as `ORA_HASH('some_string', 4)`. Hash functions can be designed with many different properties, and their design has important implications for how they are used. Figure 16-4 below includes simple ways of describing hash functions.

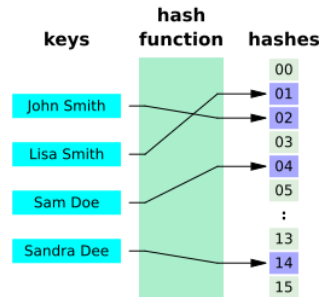
## Minimal perfect hash

Every value maps to one unique hash, and there are no empty hash buckets



## Perfect hash

Every value maps to one unique hash, but there are empty hash buckets



## Typical Hash

Multiple values may map to the same hash, and there are empty hash buckets




Figure 16-4. Description of different hash functions. Images created by [Jorge Stolfi](#) and are in the public domain.

With perfect hashing, access requires only a single read, and the operation is  $O(1)$ . With flawed hashing, where every value is mapped to the same hash, we just end up with a regular heap table stuck inside a hash bucket. In that worst-case, we have to read the entire table to find a single value, which is  $O(N)$ . There is a huge range of performance that depends on exactly how the hash is designed.

When hashing we need to be aware of space-time trade-offs. In practice we cannot achieve a minimal, perfect hash. Hash partitioning is minimal (there are no empty buckets), but far from perfect (there are many collisions) – they don't provide instant access for one row, but they don't waste much space. Hash clusters can be perfect (no collisions), but are far from minimal (many empty buckets) – they can provide instant access for a row but they waste a lot of space. Hash joins are somewhere in the middle. Those three types of hashing serve different purposes, and are described in detail in the next sections.

## Hash partitioning

Hash partitioning splits our tables into  $P$  partitions, based on the `ORA_HASH` of one or more columns. Hash partitioning will certainly not be a perfect hash; each hash bucket is a segment, which are large and meant to fit many rows. The number of hash partitions should be minimal – we don't want to waste segments.

The time to insert a row into a hash partitioned table is still  $O(1)$  – the `ORA_HASH` function can quickly determine which partition the row goes in. But the time to retrieve a row will be  $O(N/P)$ , or the number of rows divided by the number of partitions. That can be an important improvement for large data warehouse operations that read a large number of rows. But for reading a single row, that improvement is not nearly as good as what we can achieve with an  $O(\log(N))$  B-tree index access. Don't try to replace indexes with partitions, they solve different problems.

Don't try to achieve awesome hash partition performance by significantly increasing the number of partitions. If we try to build a perfect hash partitioned table, that hash will not be minimal. A large number of empty segments will waste a lot of space and cause problems with the data dictionary.

We must use the right columns for the partition key. If we use columns with a low cardinality, or don't set the number of partitions to a power-of-two, the data will not be evenly distributed among the partitions. If all of the rows are stored in the same bucket (the same hash partition) then partitioning won't help us at all.

Hash partitioning is frequently over-used in practice. To avoid abusing that feature, we need to understand the hashing trade-offs and choose a good partition key.

## Hash clusters

Hash clusters use a hash function to retrieve a *small* number of rows, unlike coarse hash partitions that return a large number of rows. In theory, hash clusters are even better than B-tree indexes for retrieving a small number of rows. The first problem with hash clusters is that we cannot simply add one to an existing table; we have to organize the table from the beginning to use a hash cluster. The hash function tells Oracle exactly where to store and find each row – there's no need to walk an index tree and follow multiple pointers. But this is a feature where theory and practice don't perfectly align.

To get  $O(1)$  read time on hash clusters we need to create a near-perfect hash. But we have to worry about that space-time trade-off. There is no practical way to get perfect hashes without also creating a huge number of unused hash buckets. We can have a hash cluster with good performance or a hash cluster that uses a minimal amount of space; we can't have both.

When we create a hash cluster we can specify the number of buckets with the `HASHKEYS` clause. In my experience, getting  $O(1)$  read time requires so many hash buckets that the table size will triple.

Unfortunately, even after all that effort to get  $O(1)$  access time, hash clusters still end up being slower than indexes. I can't explain why hash clusters are slower, this is where our theory breaks down in real-world applications.

Oracle measures the number of reads performed by a SQL statement using a statistic called "consistent gets". It's possible to create test cases where hash lookups require only 1 consistent get, while an index access on the same table requires 4 consistent gets<sup>2</sup>. But the index is still faster. For these operations, the  $O(\log(N))$  of an index is less than the  $O(1)$  of a hash cluster.

When comparing small numbers, such as 1 versus 4, the constants become more important than the Big O analysis. Hash clusters are rarely used, and Oracle surely invests more time optimizing indexes than clusters. Perhaps that optimization can compensate for the small difference between the run time complexities.

Algorithm analysis helps us drill down to the problem, but in this case the real performance difference is hidden by constants in closed source software. Perhaps this an opportunity for a custom index type. Other databases have hash indexes that don't require re-organizing the entire table to use them. Maybe someday Oracle will add that feature and make hash access work better. For now, we should ignore hash clusters.

## Hash joins

Hash partitioning can run in  $O(1)$ , with a little configuration, but only for coarse table access. Hash clusters can run in  $O(1)$ , with a lot of configuration and space, but in practice that  $O(1)$  is worse than an index's  $O(\log(N))$ . But what about hash joins?

---

<sup>2</sup> See this Stack Overflow answer, where I try and fail to create a useful constant-time index using hash clusters: <https://stackoverflow.com/questions/32071259/constant-time-index-for-string-column-on-oracle-database>

The exact time complexity is hard to pin down, but we can tell that hash joins are the best operation for joining large amounts of data.

Hashing assigns values to arbitrary hash buckets, as visualized in Figure 16-4 above. If we hash all the rows of two tables by their join columns, the matching rows will end up in the same bucket, thus joining the two tables. (The performance of hash joins will be compared with other join methods in a later section.)

Hash joins can only be used for equality conditions. The hash values are just a mapping of input values, and the relationships between the input values is broken after hashing. Input value A may be smaller than input value B, but that relationship is not necessarily true of their hash values.

Hash joins are so useful that it is worth going out of our way to enable them. We can enable hash joins by rewriting a simple non-equality condition into a weird equality condition. For example, `COLUMN1 = COLUMN2 OR (COLUMN1 IS NULL AND COLUMN2 IS NULL)` is a straight-forward, logical condition. We may be able to significantly improve the performance of joining two large tables by re-writing the condition to `NVL(COLUMN1, 'fake value') = NVL(COLUMN2, 'fake value')`. Adding logically unnecessary functions is not ideal, and may introduce new problems from bad cardinality estimates, but it's often worth the trouble if it enables faster join operations.

## Other

Constant time operations show up frequently in Oracle, like in any system. For example, inserting into a table, creating objects, and altering objects often take the same amount of time regardless of the input size.

On the other hand, all of the above-mentioned operations also have a non-constant time version. Inserting into a table can take more than constant time, if the table is index organized. Creating objects like indexes can take a huge amount of time to sort the data. Even altering tables may-or-may not take constant time, depending on the `ALTER` command. Adding a constraint usually requires validation against the data, which depends on the number of rows; but adding a default value can be done purely in metadata and can finish almost instantly.

We can't categorize operations based purely on their command type. To estimate the run time of any command we always need to think about the algorithms, data structures, and data.

## $O(\log(N))$ – Index Access

$O(\log(N))$  is the worst-case run time for index access. Index reads are another core database operation. Indexes were described in detail in Chapter 9, but a brief summary is included in the below paragraph.

When we search a binary tree, each step can eliminate half the rows of the index. Doubling the number of rows grows exponentially; conversely, halving the number of rows shrinks logarithmically. Index access on a simple binary tree is  $O(\log_2(N))$ . Oracle B-tree indexes store much more than one value per branch, and have a worst-case access time of  $O(\log(N))$ . Figure 16-5 below shows an example of a binary tree search.




Figure 16-5. Binary search algorithm. Image created by [Chris Martin](#) and is in the public domain.

We know how indexes work, and we know that indexes are great for performance, but comparing the algorithms helps us understand precisely how awesome indexes can be. Figure 16-6 below compares the fast  $O(\log(N))$  of an index read with the slow  $O(N)$  of a full table scan. This visualization is a powerful way of thinking about index performance.




Figure 16-6. Compare  $O(\text{LOG}(N))$  index read versus  $O(N)$  full table scan.

Most performance tests only compare run times at a single point along a continuum of input sizes. Thinking about the lines and curves in the above visualizations helps us understand how systems will scale. The amount of work required to read one row from an index rarely increases, no matter how large the table grows. In practice, we will never see an Oracle B-tree index grow to a height of 10, unless the data is enormous or there are strange DML patterns that cause index imbalances. Whereas full table scans grow and get slower with each new row. Indexes may not help much when our systems are small but they can make a huge difference when the system grows.

So far we've mostly discussed small operations – one hash lookup, one index lookup, or one table scan. The performance comparisons will soon become trickier as we start iterating the operations.

## $1/((1-P)+P/N)$ – Amdahl's Law

For optimal parallel processing in a data warehouse we have to worry about *all* operations. It is not enough to only focus on the biggest tasks. If our database has  $N$  cores, and we want to run large jobs  $N$  times faster, we need to parallelize *everything*.



Amdahl's law is the mathematical version of the above paragraph. The law can be expressed as the equation below in Figure 16-7<sup>3</sup>.

$$TotalSpeedup = \frac{1}{(1 - ParallelPortion) + \frac{ParallelPortion}{ParallelSpeedup}}$$

*Figure 16-7. Amdahl's law as an equation.*

We don't need to remember the equation it's only listed for completeness. But we do need to remember the lesson of Amdahl's law, which can be learned from the graph in Figure 16-8 below.

---

<sup>3</sup> Amdahl's law is not a worst-case run time complexity. But the function is important for understanding performance and worth discussing here.

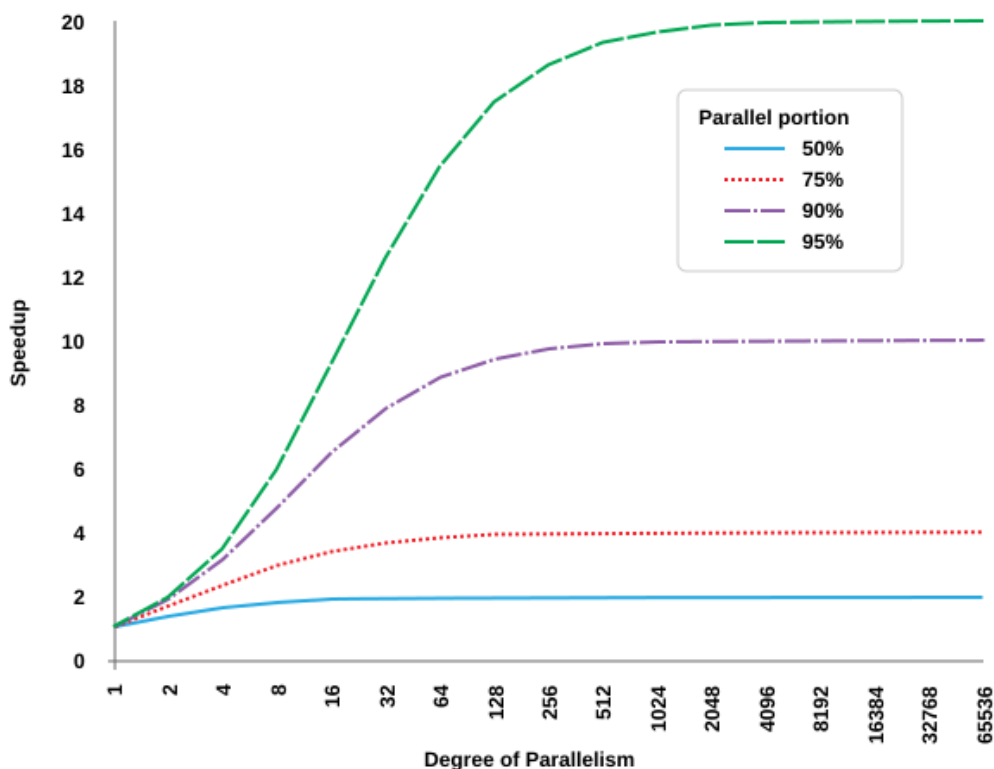


Figure 16-8. Amdahl's law as a graph. Based on ["Amdahl's Law"](#) by Daniels220, licensed under [CC BY-SA](#).

Notice that the above graph uses a logarithmic X axis, and the numbers on the Y axis are comparatively small. The implications of this graph are depressing; even the tiniest amount of serialization can dash our dreams of running  $N$  times faster.

For example, let's say we've got a large data load and we've parallelized 95% of the process. Our machine has 128 cores, and (miraculously) the parallel portion runs 128 times faster. Yet the overall process is only 17.4 times faster. If we foolishly throw hardware at the problem, and increase the cores from 128 to 256, the process will only run 18.6 times faster. Those are disappointing diminishing returns.

Our time would be much better spent increasing the parallel portion from 95% to 96%, which would increase the performance from 17.4 times faster to 21 times faster. These surprising performance gains are why we need to spend so much effort finding exactly where our processes spend their time. We all like to tune with our gut feelings, but our intuition can't tell the difference between 95% and 96%.

That level of precision requires advanced tools and a deep analysis of the process activity.

For data warehouse operations we cannot only worry about the longest-running tasks. It's not enough to parallelize only the `INSERT` or `CREATE TABLE` statements. To get the maximum speedup we need to put a lot of effort into all the other data loading operations. We need to pay attention to operations like rebuilding indexes, re-enabling constraints, gathering statistics, etc. Luckily, Oracle provides ways to run all of those operations in parallel. Previous chapters included examples of rebuilding indexes and re-enabling constraints in parallel, and Chapter 18 shows how to easily gather statistics in parallel.

To fully optimize large data warehouse operations it can help to create a graph of the system activity, with either Oracle Enterprise Manager or `DBMS_SQLTUNE.REPORT_SQL_MONITOR`<sup>4</sup>. As discussed in Chapter 12, optimizing data warehouse operations requires worrying about even the tiniest amount of system inactivity.

## O(N) – Full Table Scans, Other Operations

$O(N)$  is a simple, linear performance growth. We find this run time complexity all over the place: full table scans, fast full index scans (reading all the index leaves instead of traversing the tree), SQL parsing, basic compression, writing or changing data, etc. There's not much to see here, but things get interesting in the next section when we start comparing more functions.

Most Oracle operations fall into the  $O(N)$  category. In practice, we spend most of our tuning time worrying about the constants and hoping for a linear improvement. For example, direct-path writes can shrink the amount of data written by eliminating REDO and UNDO data. That's only a linear improvement, but a significant one.

---

<sup>4</sup> If we set the `SQL_ID` parameter to the top-level PL/SQL statement, `DBMS_SQLTUNE.REPORT_SQL_MONITOR` will generate an activity graph for all child SQL statements.

# O(N\*LOG(N)) – Full Table Scan versus Index, Sorting, Joining, Global versus Local Index, Gathering Statistics

$O(N * \log(N))$  is the worst-case run time for realistic sorting algorithms. This time complexity can show up in many unexpected places, like constraint validation. Previous sections discussed algorithms and data structures, but now it's time to start iterating those algorithms. An  $O(\log(N))$  index access is obviously faster than an  $O(N)$  full table scan, but what happens when we're looking up more than a single row?

Figure 16-9 below compares the functions discussed in this section:  $N^2$ , variations of  $N * \log(N)$ , and  $N$ .




Figure 16-9. Comparing  $N^2$ ,  $N * \log(N)$ , and  $N$ .

The above lines and curves can be used to understand full table scans versus index access, sorting, joining, global versus local indexes, gathering statistics, and many other operations.

## Full table scan versus index

Indexes are awesome for looking up a single row in a table – it's hard to beat  $O(\log(N))$  read time. But other than primary key or unique key indexes, most index access will retrieve more than one row. To retrieve multiple rows Oracle must repeat  $\log(N)$  multiple times.

If an index scan is used to lookup every row in the table, Oracle has to walk the tree  $N$  times, which leads to a  $N \cdot \log(N)$  run time. That run time is clearly much worse than just reading the entire table one-row-at-a-time. The difference is obvious if you look at Figure 16-8 above, and compare the slower, solid line of  $N \cdot \log(N)$  against the faster, dotted line of  $N$ . There are many times when a full table scan is faster than an index.

The performance difference is more complicated when Oracle reads a percentage of rows other than 1% or 100%. Notice the dashed lines in Figure 16-8 above; they represent repeating the  $\log(N)$  access for 25%, 50%, or 75% of the rows in the table. As the size of the table grows, each curve will eventually over-take the linear line for  $N$ . We cannot simply say that one algorithm is faster than another. The fastest algorithm depends on the size of the table and the percentage of rows accessed. An index might be a good idea today for reading 5% of a table, but a bad idea tomorrow if the table has grown.

There are other factors that can significantly change the balance. Oracle can use multi-block reads for full table scans, as opposed to single-block reads for index access. Using multi-block reads, Oracle can read several blocks from a table in the same time it takes to read one block from an index. And if the index clustering factor is high (if the index is not ordered by the value being searched for), an index lookup on a small number of rows may end up reading all of the blocks of the table anyway.

The theory tells us the *shape* of the lines, but only practice can tell us the actual values. I can't give you an exact number for your system, table, and workload. But there *is* a number – a point where a full table scan becomes cheaper than an index. Finding that number is not trivial, and it should not surprise us that Oracle doesn't always make the right decision.

When Oracle fails to make the right choice, we shouldn't throw out the entire decision making process by using an index hint. Instead, we should look for ways to

provide more accurate information, to help Oracle make better choices. Helping Oracle is usually done by gathering optimizer statistics.

## Sorting

$O(N * \log(N))$  is the worst-case run time for sorting. We always hope our processes will scale linearly but that is often not the case. In practice, we have to deal with algorithms that get slower faster than we anticipate. Sorting is a central part of any database, and affects `ORDER BY`, analytic functions, joining, grouping, finding distinct values, set operations, etc. We need to learn how to deal with these slow algorithms, and avoid them when possible.

For planning, we need to be aware of how performance will change over time. If sorting one million rows takes 1 second today, we cannot assume that sorting two million rows will take 2 seconds tomorrow.

We also need to be aware of how the space requirements for sorting will change with the input size. Luckily, the amount of space required grows linearly. If the number of rows doubles, the amount of PGA memory or temporary tablespace required to sort will also double.

There are times when the sorting is already done for us. A B-tree index has already been sorted – the work was done during the original `INSERT` or `UPDATE`. Adding a row to a table is  $O(1)$  – the process only requires adding a row to the end of a dumb heap. Adding a row to an index is  $O(\log(N))$  – the process needs to walk the tree to find where to add or update the value.

Oracle can use an index full scan to read from the index, in order, without having to do any sorting. The work has already been done, Oracle just needs to read the index from left-to-right or right-to-left.

Oracle can also use a min/max read to quickly find the minimum or maximum value. The minimum or maximum value in a B-tree will be either all the way on the left, or all the way on the right. Once again, the data is already sorted, finding the top or bottom result is a trivial  $O(\log(N))$  operation.

Oddly, there's a missing feature where Oracle can't find *both* the min and max using a simple min/max read<sup>5</sup>. But the below code shows a simple work-around to this problem: break the problem into two separate queries, and then combine the

---

<sup>5</sup> See my Stack Overflow answer here for more details:

<https://stackoverflow.com/q/43131204/409172>.

results. Writing an extra subquery is annoying, but it's a small price to pay for a huge improvement in run time complexity:  $O(2 * \text{LOG}(N))$  is much better than  $O(N)$ .

When we understand the algorithms and data structures used by Oracle, we know what to expect, and when to look for a faster workaround.

```
--Create a table and query for min and max values.
create table min_max(a number primary key);

--Full table scan or index fast full scan - O(N).
select min(a), max(a) from min_max;

--Two min/max index accesses - O(2*LOG(N)).
select
  (select min(a) from min_max) min,
  (select max(a) from min_max) max
from dual;
```

The set operations `INTERSECT`, `MINUS`, and `UNION` require sorting. We should use `UNION ALL` when possible because it is the only set operation that does not need to sort the results.

Sorting and joining seem to go together, but in practice they are a bad combination. The next section discusses why we don't want to use sorting for our join operations.

## Joining

Hopefully you remember the Venn diagrams and join diagrams from Chapter 1, which explained how joins logically work. Unfortunately, understanding how joins physically work is even more complicated. Figure 16-10 below visualizes the main join algorithms and their run time complexity. Each join algorithm is also described in a separate paragraph after the diagram. You may need to flip back and forth a few times to understand the algorithms. The diagram visualizes joins as the process of matching rows between two unordered lists.

# Join Algorithms and Time Complexity

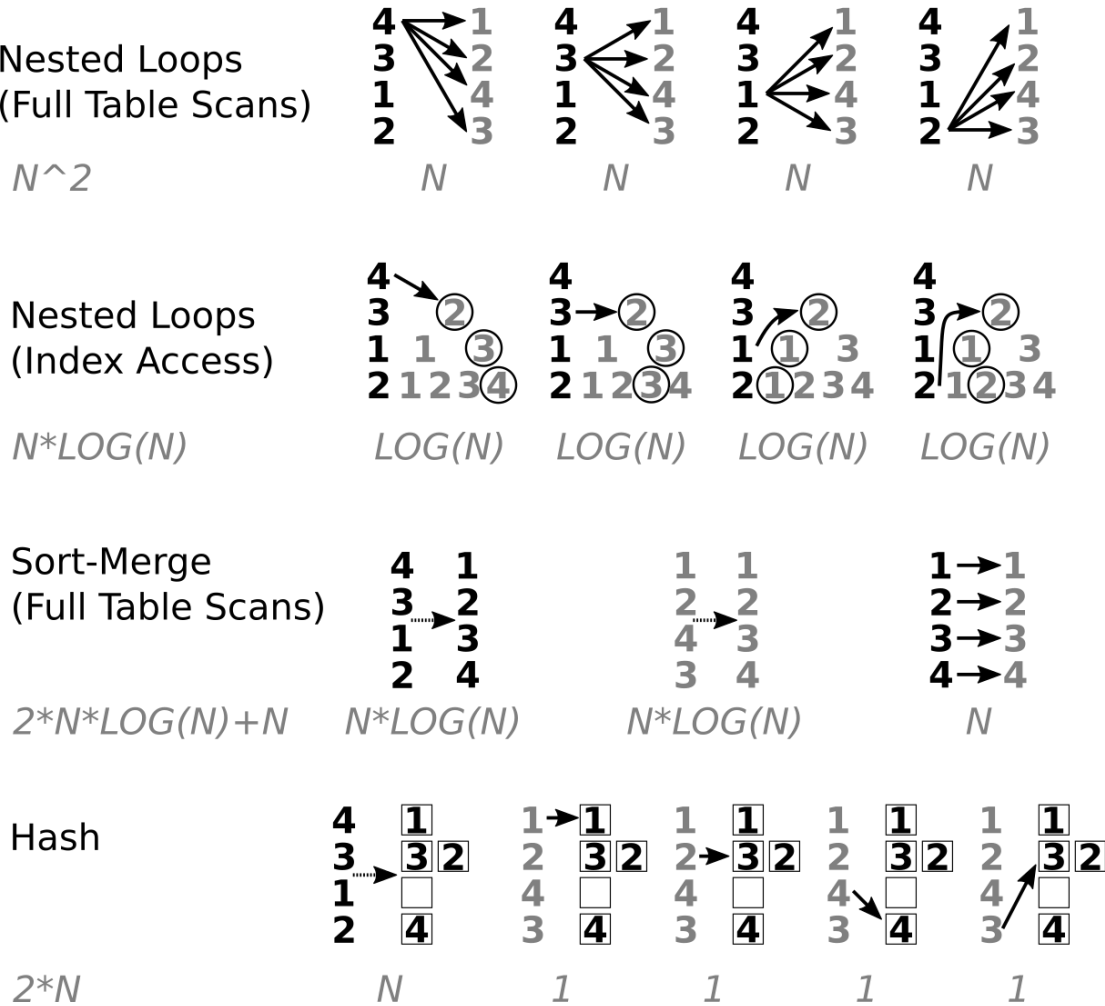


Figure 16-10. Visualization of join algorithms and time complexity.

A **nested loop with full table scans** is conceptually simple. Start with the first table, and compare each row in that table with every row in the second table. We can think of this algorithm as two nested `FOR` loops. Each phase of the algorithm does  $N$  comparisons, and there are  $N$  phases, so the run time complexity is a horribly slow  $O(N^2)$ . Recall from Figure 16-9 that the line for  $N^2$  looks almost completely vertical. This join algorithm should only be used when the tables are trivially small, the optimizer statistics are bad, or important indexes are missing.

A **nested loop with index access** is a lot smarter. Instead of searching an entire table for each row, we can search an index. Reducing  $O(N^2)$  to



$O(N * \log(N))$  is a huge improvement. To better understand this run time, it might help to use separate variables for the two table sizes; let's divide  $N$  into  $M$  and  $N$ . If we think about  $O(M * \log(N))$ , the shape of the curve mostly depends on  $M$  instead of  $N$ . Even if one of the tables is huge, if we only have to do a small number of index accesses, the performance will be great. But if both  $M$  and  $N$  are large, this join algorithm is not efficient. Nested loops operations work best when the number of rows is small for one table, and there is index access into the other table.

A **sort-merge join with full table scans** operates by first sorting both tables, which is expensive, and then matching the sorted rows, which is cheap. In practice we don't see sort-merge joins very often – nested loops with index access is a better choice for small joins, and a hash join is a better choice for large joins. But if there are no relevant indexes (ruling out nested loops), and there is no equality condition (ruling out a hash join), then sort-merge is the best remaining option. If one of the tables is already presorted in an index, then half of the sorting work is already done.

A **hash join** has two phases. The smaller table is read and built into a hash table, ideally in memory. Then the larger table is scanned, and the hash table is probed for each row in the larger table. With a perfect hash, writing and reading from a hash table only takes one step, and the run time is only  $O(2 * N)$ . But in practice there will be collisions, and multiple rows will be stored in the same hash bucket. The time to perform hash joins does not grow linearly; hash joins are somewhere between  $O(2 * N)$  and  $O(N * \log(N))$ <sup>6</sup>.

Hash joins have a great run time complexity but that doesn't mean we always want to use them. Hash joins read all of the rows of both tables, even if we are only going to match a few rows. Hash joins require a lot of memory or temporary tablespace, roughly equal to the size of the smaller input table. A slow algorithm in memory might be better than a fast algorithm on disk. And hash joins are not always available, because they only work with equality conditions.

There are many variations of the above join algorithms. Cross joins, also known as Cartesian products, are similar to nested loops with full table scans. Parallelism and partitioning can always complicate the algorithms. Hash joins can use bloom filters to eliminate values without fully comparing all of them. Joining is the most important operation in the database, there's not enough space here to describe all of the features for joining tables.

---

<sup>6</sup> See my answer here for tests comparing sort-merge with hash: <https://stackoverflow.com/a/8548454/409172>. When available, hash joins are always faster.

Join performance is all about Oracle correctly estimating the size of tables. The table size determines which algorithms are used, and how the tables are used in the algorithms.

Hashing and sorting are bad ideas if one of the tables returns no rows – there's no need for all that prep work when the intermediate data structures won't be used. A nested loop is a horrible idea for two huge tables that have a lot of matching rows. Bad execution plans happen when Oracle thinks a small table is large, or a large table is small. How that mistake happens, and how to fix it, is discussed in Chapters 17 and 18.

## Global versus local indexes

The advantages and disadvantages of partitioning are different for tables and indexes. Partitioning a table is often a good choice – there are many potential performance improvements and few costs. Indexes can be either global, one index for the whole table, or local, one index per partition. Compared to table partitioning, partitioned index benefits are smaller, and the costs are greater.

Reading from a single table partition, instead of a full table scan, significantly decreases the run time complexity from  $O(N)$  to  $O(N/P)$ , where  $P$  is the number of partitions. Reading from a local index, instead of a global index, only changes the run time complexity from  $O(\log(N))$  to  $O(\log(N/P))$ . That index access improvement is barely noticeable. (Unless the data is skewed differently between partitions, and some partitions have indexes that are more efficient than others.)

The cost of reading an entire partitioned table, without partition pruning, is still a normal  $O(N)$ . But reading from a local index without partition pruning is much more expensive than a global index read. Reading from one large index is  $O(\log(N))$ . Reading from many small indexes is  $O(P \cdot \log(N/P))$ , a significant increase. Walking one big tree is much faster than walking multiple small trees. We need to think carefully when we partition, and we should not expect table and index partitioning to work the same way.

## Gathering Optimizer Statistics

Gathering optimizer statistics is an essential task for achieving good SQL performance. Statistics are used to estimate cardinality, the number of rows returned by an operation. Cardinality is synonymous with “Input Size”, the X axis on

most of the graphs in this chapter. Cardinality is vital information for making execution plan decisions, such as knowing when to use a nested loop or a hash join.

To get good execution plans we need to gather statistics periodically and after any significant data changes. But gathering statistics can significantly slow down our process if we're not careful. Understanding the different algorithms and data structures for gathering statistics can help us avoid performance problems.

Finding the cardinality isn't just about counting the absolute number of rows. Oracle also needs to measure the distinctness of columns and specific column values. For example, an equality condition on a primary key column will return at most one row, and is a good candidate for an index access. But an equality condition on a repetitive status column is more complicated; some statuses may be rare and benefit from index access, other statuses may be common and work best with a full table scan. Optimizer statistics gathering needs to generate results that can answer those more complicated questions.

Counting distinct items is similar to sorting, which is a slow operation. To make things even worse, Oracle needs statistics for all the columns in a table, which may require multiple passes. A naïve algorithm to count distinct values would first sort those values; it's easy to measure distinctness if the values are in order. But that naïve approach would take  $O(N \cdot \log(N))$  time. A better algorithm would use hashing, and that's exactly what recent versions of Oracle can do with the `HASH GROUP BY` and `HASH UNIQUE` operations. As with joining, the time to hash is somewhere between  $O(N)$  and  $O(N \cdot \log(N))$ .

Luckily, when gathering optimizer statistics we can trade accuracy for time. The default Oracle statistics gathering algorithm performs an *approximate* count with a single pass of the table. The single pass algorithm runs in  $O(N)$  time and generates numbers that are accurate but not perfect.

Developers rightfully worry about imperfect values. A single wrong bit can break everything. But in this case an approximate value is good enough. The optimizer doesn't need to know the *exact* value. Oracle only needs to know if the algorithm needs to be optimized for "large" or "small" queries.

The example below shows how close the approximation is. This code uses the new 12c function `APPROX_COUNT_DISTINCT`, which uses the same algorithm as statistics gathering.

```
--Compare APPROX_COUNT_DISTINCT with a regular COUNT.
select
    approx_count_distinct(launch_date) approx_distinct,
    count(distinct launch_date)         exact_distinct
from launch;
```

| APPROX_DISTINCT | EXACT_DISTINCT |
|-----------------|----------------|
| -----           | -----          |
| 61745           | 60401          |

The fast approximation algorithm<sup>7</sup> only works if we read the entire table. If we try to sample a small part of the table, the approximation algorithm no longer applies, and Oracle has to use a different approach. We may think we're clever setting a small estimate percent, such as:

```
DBMS_STATS.GATHER_TABLE_STATS(..., ESTIMATE_PERCENT => 50).
```

But a slow algorithm reading 50% of the table is slower and less accurate than the fast algorithm reading 100% of the table. We should rarely, if ever, change the `ESTIMATE_PERCENT` parameter.

These distinct counting tricks can also apply to partition statistics, using a feature called incremental statistics. Normally, partitioned tables require reading the table data twice to gather statistics; one pass for each partition, and another pass for the entire table. The double read may seem excessive at first, but consider that we cannot simply add distinct counts together.

But incremental statistics uses an approximation algorithm that does enable adding distinct counts. Incremental statistics creates small data structures called synopses, which contain information about the distinctness within a partition. After gathering statistics for each partition, the global statistics can be inferred by merging those small synopses. This algorithm improves statistics gathering from  $O(2*N)$  to  $O(N)$ . That decrease may not sound impressive, but remember that partitioned tables are often huge. Saving an extra full table scan on our largest tables can be a big deal.

This algorithm analysis is starting to get recursively ridiculous. We're discussing algorithms that help us determine which algorithms to use. And the problem goes deeper – in rare cases, statistics gathering is slow because the optimizer chooses a bad plan for the statistics gathering query. For those rare cases we may need to prime the pump; we can use `DBMS_STATS.SET_*_STATS` to create initial, fake statistics, to help us gather the real statistics.

Gathering statistics may feel like slow, annoying maintenance work. But instead of ignoring optimizer statistics, we should try to understand the underlying algorithms and use them to our advantage to improve the run time of our systems.

---

<sup>7</sup> Oracle uses an algorithm called HyperLogLog for distinct count approximations.

# $O(N^2)$ – Cross Joins, Nested Loops, Other Operations

$O(N^2)$  algorithms are getting into the ridiculously-slow territory. Unless we have special conditions, such as an input size close to 0, we should avoid these run times at all costs. This run time complexity happens from cross joins, nested `FOR` loops, nested loop joins with full table scans, and some other operations. As a quick reminder of their poor performance, Figure 16-11 below shows  $N!$  and  $N^2$ .

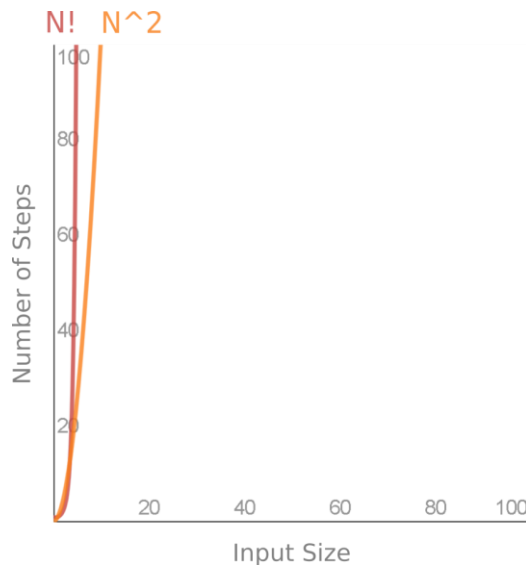


Figure 16-11. Compare  $N!$  with  $N^2$ .

Cross joins are not always bad. Sometimes it is necessary to mindlessly combine every row in one table with every row in another table. And cross joins can be fast if the number of rows is close to zero. But cross joins that happen because of incorrectly joined tables, or because of poor optimizer cardinality estimates, are horrible for performance. Accidental cross joins are so bad that they can effectively bring down an entire database. Oracle may need an unlimited amount of temporary tablespace to build the result set, depriving other processes of space for sorting or hashing.

`FOR` loops inside `FOR` loops are the most common way to generate  $O(N^2)$  run times in procedural code. If we add *another* `FOR` loop the run time complexity becomes  $O(N^3)$ , then  $O(N^4)$ , etc. Nested `FOR` loops can easily happen in PL/SQL, and they are all-too-common when using explicit cursors. Avoiding that

horrible performance is an important reason why we should always try to replace nested cursors with a single SQL statement.

Developers write procedural code as nested `FOR` loops because loops are the easiest way to think about joins. With imperative code, Oracle cannot fix the algorithms for us – Oracle is stuck doing exactly what we ask for. The advantage of switching to declarative code is that we don't have to worry about these algorithms as much. When we use SQL we let Oracle decide which join algorithms to use.

It is possible that after switching to declarative SQL the optimizer will still choose an  $O(N^2)$  algorithm. As we saw in Figure 16-10 earlier in the chapter, a nested loops join with two full table scans is a terrible way to join tables. But that worst-case should only happen if we have missing statistics, missing indexes, or weird join conditions.

$O(N^2)$  shows up in other unexpected places. In extreme cases the SQL parse time seems to grow exponentially. For example, the parse time becomes ridiculously bad if we combine thousands of subqueries with `UNION ALL`. (The code measuring parse time is too large to include here but it can be found in the repository.) Writing large SQL statements is a good thing but writing Brobdingnagian SQL is not.

The `MODEL` clause brings procedural logic to SQL statements. `MODEL` is a neat trick and explained briefly in Chapter 19. That clause gives us the ability to turn our data into a spreadsheet and apply `FOR` loops to our rows and columns. Just like with PL/SQL, we can easily find ourselves in  $O(N^2)$  performance territory.

A high run time complexity isn't always avoidable or necessarily bad. But we should think carefully when we see cross joins or nested operations.

## O(N!) – Join Order

$O(N!)$  is as bad as Oracle's algorithms get, but luckily it's also rare. As discussed in Chapter 6, tables have to be joined in a specific order, and there are many ways to order a set of tables. The problem of ordering tables must be solved by the optimizer when building execution plans, and in our minds when trying to understand queries.

The optimizer is capable of handling a huge number of tables without serious problems. But there are always unexpected problems when we push the system to the limits. For example, if we nest over a dozen common table expressions the parse time appears to grow like  $O(N!)$ . (See the GitHub repository for a code demonstration.) In practice, this problem will only happen because of a rare bug or because we're doing something we shouldn't.

Our minds can only hold a small ordered list in short term memory. If we build our SQL statements with giant comma-separated lists we'll never be able to understand them. Building small inline views and combining them with the ANSI join syntax will vastly simplify our SQL. There's no precise equation for SQL complexity, but we can think of the complexity comparison like this:  $(1/2N)! + (1/2N)! \ll N!$

For performance we want to batch things together, to reduce overhead. But to make our code readable we want to split everything up into manageable pieces. That is precisely what SQL excels at; we can logically divide our queries to make them understandable, and then the optimizer can efficiently put the code back together.

## $O(\infty)$ – The Optimizer

The Oracle optimizer builds SQL execution plans. Building optimal execution plans is impossible, so we could say this task is  $O(\infty)$ . That run time complexity sounds wrong at first; Oracle obviously builds at least *some* optimal execution plans. But it's helpful to think of the optimizer as doing an impossible task.

Building the best execution plan means the optimizer must compare algorithms and data structures, and determine which one runs fastest. But it is literally impossible to determine if a generic program will even *finish*, much less how long it will run. Determining if a program will finish is called the halting problem. Before computers were even invented, Alan Turing proved that it is impossible for computers to solve the halting problem. Luckily, this is a problem that's theoretically impossible to solve in all cases, but is practically feasible to solve in almost all cases.

To add another wrinkle to the problem, Oracle has to generate the execution plan incredibly fast. A better name for the optimizer would be the satisficer. Satisficing is solving an optimization problem while also taking into account the time necessary to solve the problem.

It's important we understand how difficult the optimizer's job is. We shouldn't be surprised when Oracle occasionally generates a bad execution plan. The optimizer isn't as bad as we think, it's merely trying to solve an unsolvable problem. When the optimizer is having trouble we shouldn't abandon it, we should try to work with it.

When a bad execution plan is generated our first instinct shouldn't be, "how can I work around this bad plan?" Our first instinct should be, "how can I provide

better information so Oracle can make better decisions?” That information is almost always in the form of optimizer statistics, which are described in Chapters 17 and 18.

We need to resist the urge to quickly use hints and change system parameters. Think of the optimizer as a complex forecasting system. We should never change a system parameter such as `OPTIMIZER_INDEX_COST_ADJ` just because it helps with one query. That would be like adding 10 degrees to every weather forecast because the meteorologist was once wrong by 10 degrees.

## Summary

Many performance problems fall into a small number of run time complexities. Knowing which functions represent our problems can help us understand why Oracle is behaving a certain way, and how to find a solution. Practical algorithm analysis is simply matching our problems with pre-defined classes of problems. We may not use this approach often, but without it we'll never be able to truly understand database performance.

Most of our performance problems are related to implementation details, and those pesky constants we've been conveniently ignoring. The next chapter looks at a more traditional list of Oracle performance concepts.



**You are viewing a courtesy copy of chapter  
16 from Jon Heller's book, *Pro Oracle SQL  
Development***

**Would you like a copy of the complete  
book?**

**Purchase now from Amazon:**

**[https://www.amazon.com/Pro-Oracle-SQL-  
Development-Practices/dp/1484245164/](https://www.amazon.com/Pro-Oracle-SQL-Development-Practices/dp/1484245164/)**

**Or from Apress.com:**

**[https://www.apress.com/us/book/97814842  
45163](https://www.apress.com/us/book/9781484245163)**

**Thank you for your support of the author  
and his work**